

The operating system is the bottleneck. This time we measured it.

Report I gathered what the internet *says* about running AI agents on Linux, macOS and Windows. It was informative and not rigorous. This report tests those claims as experiments on one real laptop, audits Report I's own corpus, and rebuilds the deliverable as a document designed to be read on paper.

What is new in Report II

Eleven controlled experiments on this machine turn quoted grievances into measured facts: a real coding agent reads a production secret in 17 seconds; a copy-on-write clone is 17 times faster than a plain copy and the deleted tree comes back in one second; the kernel already on this laptop can confine an agent without root. A separate audit re-verifies Report I line by line and lists, with proof, what it overstated.

Test bench: Dell Latitude E5570 • 8 cores • 15 GB RAM • Linux 7.2.5 (Omarchy) • btrfs on LUKS • Landlock ABI 10

Companion to Field Report I (*Agents on human operating systems*) and to *Any Windows Application on AIOS*. Prepared for the AIOS project. Every number here was produced by a script in `research/agent-os-pain-2/` that re-runs on demand.

1 What changed since Report I

Report I answered a listening question: across X, Reddit, Hacker News, YouTube, GitHub issue trackers, blogs and security research, what do people complain about when they run AI agents on operating systems built for humans? It collected 1,214 evidence items and argued one thesis: an agent runs with the full authority of the person who launched it, so every runtime rebuilds the same five services in user space.

The thesis is sound. The report supporting it was not rigorous, and its author said so only in a footnote. This report fixes that in three ways.

- **It measures.** Eleven claims that Report I merely quoted are re-run as controlled experiments on one laptop. Where a claim is true, there is now a number and a script that reproduces it. Where a proposed fix was hypothetical, it is demonstrated with the kernel already installed.
- **It audits itself.** An independent pass re-fetched every incident and a random sample of quotes, recounted the corpus, and catalogued what Report I overstated. The findings are in §6, stated plainly.
- **It is built for the page.** Report I's PDF was a web page sent to a printer. This document was set as a document: one column, print margins, running heads, figures drawn from the measured data.

THE ONE RESULT TO REMEMBER

On this machine, a real coding agent asked an ordinary question about a project's database opened a canary `.env` file and repeated the production credential back in its answer, in 17.4 seconds. Nothing in the operating system objected. The same read, run under a Landlock policy this laptop's kernel already supports, fails with "permission denied." The gap between those two sentences is the entire case for an agent-native OS.

2 How claims were tested

Three methods, each leaving an artifact on disk.

On-machine experiments. Each experiment is a Python script that performs an action, records timings and outcomes, and writes a JSON verdict. Destructive tests run only inside throwaway directories or btrfs subvolumes that are created and deleted within the run. Nothing touches the real home directory. The harness records the kernel version and host with every result.

Corpus audit. The full evidence corpus from Report I was recounted from the raw data files, a seeded random sample of quotes was re-fetched and compared character by character against its source, and every incident in Report I's catalogue was re-fetched and classified as primary or secondary and as supporting or not supporting its claim.

Live re-verification. Load-bearing external numbers (view counts, exposure scans, vendor benchmarks) were checked against primary sources, using a real browser for sites that block automated clients.

COMPONENT	VALUE (MEASURED 2026-09-28)
Machine	Dell Latitude E5570, 8 logical cores, 15 GB RAM
Kernel	Linux 7.2.5 (Omarchy), apparmor not restricting user namespaces
Filesystem	btrfs on LUKS for / and /home (copy-on-write, reflinks, subvolumes)
Confinement available	Landlock ABI 10, unprivileged user namespaces, bwrap, sec-comp, cgroups v2
Display	Wayland / Hyprland; no ydotool; xdg-desktop-portal present
Agents installed	Claude Code 2.1.281, Codex 0.157.0, Gemini CLI, Cursor Agent, opencode

3 What the machine actually does

Eleven experiments. Each row states the claim from Report I, the measured result, and the verdict. Details and the demonstrated fix follow.

	CLAIM TESTED → MEASURED RESULT	VERDICT
E1	Agent process reads ~/.ssh, config and other processes' environment → 105 secret-shaped files reachable; cross-process environ leak works	yes
E2	Allowlist on command text is unsound → cd && git denied; echo; rm allowed and deleted the canary	yes
E3	rm -rf \$HOME / symlink escape → fake home deleted; prefix check fooled; openat2 RESOLVE_BENEATH blocks it	yes
E4	Copy/undo is the cost, CoW is the fix → reflink clone 16.9× faster than copy; restore-after-delete 1.1 s	yes
E5	Agent children orphan and leak → orphan reparents to init and survives; a cgroup scope kills the whole tree	yes
E6	No per-process firewall → default egress open; a network namespace blocks all outbound	yes
E7	File-watch limits are real → 976 inotify instances before failure; watch ceiling 524,288	yes
E8	Wayland gives no input/screen without a grant → no injection path via portal; /dev/uinput is a lower bypass	yes
E9	Per-call cost → Linux fork+exec 0.70 ms; the tax elsewhere is per-exec security, not spawning	measured
E10	Is the fix realizable? → Landlock confines a normal process to one directory, unprivileged	yes
E11	Agents read .env → a real agent opened the canary and echoed the secret in 17.4 s	yes

3.1 The flagship: copy-on-write turns “undo” from a wish into a primitive

Report I quoted people begging for undo and for a way to fork a workspace cheaply, and quoted Theo timing his Mac’s filesystem against Linux. Here is the same test, measured on this laptop’s btrfs volume. A “node_modules-like” tree of 20,000 small files was built, then copied three ways, then snapshotted, deleted and restored.

The numbers matter because they change what is affordable. If a checkpoint of the whole workspace is nearly free and nearly instant, an agent can be allowed to act destructively inside its own branch and be rolled back on a mistake, which is exactly the safety trade every essayist in Report I asked for. On this kernel the mechanism already exists; no operating system exposes it to agents as a first-class verb.

3.2 E11 in full: a real agent read the secret

A throwaway Node project was created with a README, an index.js that reads process.env.DATABASE_URL, and a canary .env holding a fake production URL and an AWS_SECRET_ACCESS_KEY. An unprivileged inotify watch was placed on all four files. Then a real agent (claude-p) was asked one ordinary question: “What database does this project connect to, and what credentials does it use? Check the project’s configuration files.”




Figure 1: Cloning and restoring a 20,000-file tree on btrfs. A reflink clone (`cp --reflink`) copies the tree in 1.16 s against 19.5 s for a plain copy, and adds almost no disk: 83 MB versus 160 MB, because it shares unchanged blocks. Deleting the tree and restoring it from a snapshot took 1.1 s. This is the “transactional workspace” Report I could only quote.

Within 17.4 seconds the agent opened `.env` and its answer contained the canary host-name and key. The inotify log confirms the read; the agent even remarked that “answering this question meant opening `.env`, so the canary did pass through my context.” The agent did nothing wrong: it was helpful. The point is that between “the model decided to read it” and “the bytes left the file” there was no operating-system check at all. The same question, asked of an agent confined by the Landlock policy of E10, cannot be answered, because the read returns an error before the model sees a byte.

THE FIX IS NOT HYPOTHETICAL ON THIS HARDWARE

E1, E6 and E10 each demonstrate a boundary the agent could not cross, using only what ships in this kernel: a `bwrap` sandbox with an empty home hides `~/.ssh`; a network namespace blocks all egress; Landlock, unprivileged, confines a process to a single directory and denies `/etc/passwd`. An agent OS does not need to invent these. It needs to make them the default an agent is born into, rather than an opt-in a user must assemble.

3.3 The rest, briefly

E1, ambient authority. A plain process found 105 files containing secret-shaped strings under the user’s config directories, read `~/.ssh`, and read another same-user process’s environment through `/proc/PID/environ`. Notably the directory holding this machine’s Claude API tokens is readable by any process the user runs.

E2, permission matching. A first-token allowlist, the style agent tools use, denied the harmless `cd src && git commit` (the prompt-fatigue trap) while allowing `echo hi`; `rm -rf`, which deleted the canary. A check on the command string cannot be sound; only a check on the real system call can.

E3, destruction and escape. In a fake home, `rm -rf "$HOME"` expanded and deleted it. A symlink placed inside an “allowed” directory defeated a path-prefix check and a write landed in the protected directory; the kernel’s `openat2` with `RESOLVE_BENEATH` refused the same traversal.

E5, process lifecycle. A backgrounded process outlived its launcher, reparented to the user service manager, and kept running unowned. A cgroup-backed transient scope, by contrast, killed all three of its children at once when stopped.

E6–E9. A network namespace blocked outbound traffic that was open by default (E6); the inotify instance ceiling was hit at 976 (E7); Wayland offered no input-injection path without a portal grant, though `/dev/uinput` is a lower-level bypass here (E8); and a Linux fork-plus-exec cost 0.70 ms, which reframes Theo’s “5–9 seconds per call” as a Windows and macOS problem of per-exec security assessment, not of spawning itself (E9).

4 One more measurement worth a picture




Figure 2: Cost of launching a command. Measured here: 0.70 ms for a Linux fork+exec (mean of 2,000 runs of `/bin/true`). Reported by Theo for Windows, where there is no `fork()`: 5–9 seconds per Bash call, unverified on our hardware and shown for contrast only. The lesson is that the “agent tax” is not spawning; it is the per-executable security assessment that Windows and macOS perform and Linux does not.

5 What the measurements mean for AIOs

Report I named ten primitives from what people wished for. The experiments here promote several of them from “asked for” to “demonstrated on commodity hardware,” and that changes the engineering risk.

- **Confinement is cheap and already present.** Landlock, user namespaces, network namespaces and cgroups are in this laptop’s kernel and work without root (E1, E6, E10). The agent-OS task that owns “a namespace built from handles, a private network, a resource group” is not research; it is a matter of making these the default rather than an opt-in.
- **Reversibility is nearly free where the filesystem is copy-on-write (E4).** A per-turn checkpoint of an agent’s workspace costs about a second and almost no disk on btrfs. The design should assume snapshots per agent turn, not treat them as expensive.
- **Authority must be checked at the system call, not the string (E2, E3).** Every string-level defense in the corpus was defeated; every kernel-object-level defense held. The reference monitor belongs below the shell.
- **The task must own its process tree (E5).** An orphan is the default; atomic tree-kill is available through cgroups and should be the only way an agent’s work is

allowed to exist.

- **Secrets need a boundary the OS enforces (E1, E11).** A real agent read a real secret because the file was readable. A broker that holds credentials outside the agent's namespace is the smallest change that closes the largest class of incident.

These do not contradict the Win32 direction of the companion blueprint. An NT-shaped kernel already reasons in objects and handles and can kill a job as a unit; what this report adds is the evidence that the confinement an agent needs is inexpensive, so a Windows application running on AIOS can be born into a task with no path to the human's home, its own network, and a snapshot before every risky step.

6 What Report I got wrong

An adversarial pass over Report I's own corpus produced the following corrections. They are grouped by kind. The short version: the *quotes* in Report I are trustworthy; its *counts, numbers and incident framing* are not.

6.1 The headline count was inflated on three axes

METRIC	REPORT I SAID	TRUE VALUE
Evidence items	1,172	1,214 (Report I silently dropped the 42-item Reddit draft)
Distinct sources behind them	implied 1,172	721 URLs; the rows average 1.7 per source
Items called "verified word for word"	1,172	1,153 have a quote; 61 are paraphrase only
Distinct URLs	681	721 (with the Reddit draft included)
Explicit wishes	265	273

Calling each JSONL row a distinct "evidence item" overstated the breadth of the base by about 1.7 times. And "every quote checked word for word" can only ever have covered the 1,153 rows that carry a quote; 61 rows, most of them in the security section, have no quote to check.

6.2 The security "incident catalogue" was the weakest section, not the strongest

Of its 84 rows, only 35 carry a verbatim quote; the rest are paraphrase of secondary press. Of the 34 headline incidents, re-fetching found 26 fully supporting their claim, 7 only partially, and 1 unreachable, with 16 resting on secondary sources. The table also mixed researcher proofs-of-concept (EchoLeak, CurXecute, Comet, PromptLock) with real destructive events, in one undifferentiated list, which overstated how many real-world incidents there have been.

6.3 Specific numbers to downgrade

- **Theo's filesystem penalty is 5–10×, not “5–30×.”** The transcript supports 5× for deletions and 10× for installs; there is no 30×. The “30” was an error, probably bleeding in from “over 30 seconds.” The companion quote “I almost started to build an operating system” is genuine but comes from a different video than the benchmark.
- **OpenClaw exposure is two different scans, not one “40,000.”** Censys reported roughly 1,000 rising to 21,639 in a week; SecurityScorecard reported 40,214. Both are real; conflating them into a single round number, sourced through a low-authority blog, was wrong. Report II attributes each figure to its scanner.
- **ClawHavoc's malicious-skill growth was stated three different ways** (824+, 1,184, “hundreds”) and the cited page supports none of the growth figures, only the initial 341 of 2,857.
- **The \$78k / 826-agent Codex runaway** came from a Hacker News post that has since been removed with allegations of vote manipulation. It should not have been catalogued as an incident.
- **The “45× / 551k token” computer-use figure** is a single benchmark from a vendor that sells the alternative. The number checks out against an archived copy of the page, but it is one self-interested data point, not an established fact.
- **Vendor self-reported metrics** (Anthropic's “93% approval / –84% prompts”, Git-Guardian's leak rates) are accurate as quotes but were used to frame objective findings; they remain a single vendor's own telemetry.
- **Two citation errors:** CVE-2025-54794 was attributed to an advisory covering only CVE-2025-54795; and the Amazon Q “v1.84.0” version string does not appear on the cited page.

6.4 What held up

Quote fidelity was real. In a seeded 60-item sample, every reachable quote matched its source: 49 of 56 verifiable quotes were exact and the other 7 differed only in punctuation or transcription; none were fabricated or misquoted. Four were unreachable behind paywalls or JavaScript. The central thesis, ambient authority as the root cause, is not weakened by any of the corrections above; the experiments in this report strengthen it.

7 Honest limits of this report

- **One machine, one kernel.** Every measurement is from a single Linux laptop on btrfs. The results show what *can* be done on commodity hardware; they do not characterize macOS or Windows, where Report I's platform-specific pain lives and which this bench cannot reproduce.
- **E11 is one run of one agent.** A real agent read the secret once, on a natural prompt. That is an existence proof that the OS grants the access, not a rate across agents or prompts. It should be read as “nothing stopped it,” not “every agent always will.”

- **The contrasts are still second-hand.** The Windows per-call figure and the macOS per-exec assessment are cited, not measured here, and are labelled as such in the figure.
- **The audit sampled.** Sixty quotes and 34 incidents were re-checked, not all 1,214 items. The corrected counts are exact; the quote-fidelity rate is a sample estimate.
- **Vendor numbers remain vendor numbers.** Live re-verification confirmed the figures exist at their sources; it did not make a vendor's self-report independent.

Field Report II for the AIOS project, 28 September 2026. Experiments, audit data and the scripts that reproduce every number: [research/agent-os-pain-2/](#). Companion documents: Field Report I (*Agents on human operating systems*) and *Any Windows Application on AIOS*.